



## **cv: An R Package for Cross-Validating Regression Models**

**John Fox**   
McMaster University

**Georges Monette**   
York University

---

### **Abstract**

We describe the **cv** package, which implements cross-validation for standard R regression models through a uniform and simple to use generic function. Methods for the **cv()** function are provided for many commonly employed classes of statistical models, including mixed-effects models, and it is straightforward to extend the **cv** package by writing methods for additional classes of models. The **cv()** function also can cross-validate complex model-selection procedures, such as those that include variable transformations, predictor selection, and selection among competing models. The **cv()** function generally supports parallel computations, and the supplied methods for linear and generalized linear models take advantage of computational efficiencies.

*Keywords:* cross-validation, regression analysis, model selection, R.

---

## **1. Introduction**

Cross-validation (CV) is an essentially simple and intuitively reasonable approach to estimating the predictive accuracy of regression models. CV is developed in many standard sources on regression modeling and “machine learning”—we particularly recommend [James, Witten, Hastie, and Tibshirani \(2021, Secs. 5.1, 5.3\)](#)—and so we will describe the method only briefly here before taking up computational issues and some examples. See [Arlot and Celisse \(2010\)](#) for a wide-ranging, if technical, survey of cross-validation and related methods that emphasizes the statistical properties of CV.

Validating research by replication on independently collected data is a common scientific norm. Emulating this process in a single study by data-division is less common (see, e.g., [Barnard](#)): The data are randomly divided into two, possibly equal-size, parts; the first part is used to develop and fit a statistical model; and then the second part is used to assess the adequacy of the model fit to the first part of the data. Data-division, however, suffers from

two problems: (1) Dividing the data decreases the sample size and thus increases sampling error; and (2), even more disconcertingly, the results can vary substantially based on the random division of the data, particularly in smaller samples. See [Harrell \(2015, Sec. 5.3\)](#) for this and other remarks about data-division and cross-validation.

Cross-validation speaks to both of these issues. In CV, the data are randomly divided as equally as possible into several, say  $k$ , parts, called “folds.” The statistical model is fit  $k$  times, leaving each fold out in turn. Each fitted model is then used to predict the response variable for the cases in the omitted fold. A CV criterion (also termed a “cost” or “loss” measure), such as the mean-squared error (“MSE”) of prediction, is then computed using these predicted values. In the extreme  $k = n$ , the number of cases in the data, thus omitting individual cases and refitting the model  $n$  times—a procedure termed “leave-one-out (LOO) cross-validation.”

Because the  $n$  models are each fit to  $n - 1$  cases, LOO CV produces a nearly unbiased estimate of prediction error. The  $n$  regression models are highly statistically dependent, however, based as they are on nearly the same data, and so the resulting estimate of prediction error has relatively large variance. In contrast, estimated prediction error for  $k$ -fold CV with  $k = 5$  or 10 (commonly employed choices) are somewhat biased but have smaller variance. It is also possible to correct  $k$ -fold CV for bias (see Section 7.2).

The **cv** package for R automates the process of cross-validation for standard R statistical model objects. The principal function in the package, also named `cv()`, has methods for objects produced by a number of commonly employed regression-modeling functions, including those for mixed-effects models:

```
R> library("cv", quietly=TRUE)
R> methods("cv")
```

```
[1] cv.default*   cv.function* cv.glm*       cv.glmmTMB*   cv.lm*
[6] cv.lme*       cv.merMod*   cv.modList*  cv.rlm*
see '?methods' for accessing help and source code
```

- The “`lm`” and “`glm`” methods are for linear and generalized-linear models fit respectively by the standard R `lm()` and `glm()` functions.
- The “`modList`” method for `cv()` cross-validates several competing models, not necessarily of the same class, using the same division of the data into folds. The `cv()` function is introduced in the context of a preliminary example in Section 2 of the paper.
- Cross-validating mixed-effects models, using the “`glmmTMB`”, “`lme`”, and “`merMod`” methods for `cv()`, involves special considerations that we take up in Section 3.
- The “`function`” method for `cv()`, discussed in Section 4, cross-validates a potentially complex model-specification process that may, for example, involve choice of data transformations and predictors, or model selection via CV itself.
- The “`default`” `cv()` method works (perhaps with a bit of coaxing) with many other existing regression-model classes for which there is an `update()` method that accepts a `data` argument. More generally, the **cv** package is designed to be extensible, as discussed in Section 5.

A number of existing R packages include functions for cross-validating regression models. We briefly situate the **cv** package relative to other R software for cross-validation in Section 6.

The final section of the paper (Section 7) describes some computational details, including efficient CV computations for linear and generalized-linear models and computation of bias-corrected CV criteria.

In the interest of brevity, we won't describe all of the features of the **cv** package here, concentrating on the aspects of the package that are relatively novel. For example, the `cv()` function can perform computations in parallel and can independently replicate a cross-validation procedure several times. There are also data management facilities in the package, such as coercing the objects produced by the `cv()` function into data frames for further analysis. These and other features not discussed in this paper are taken up in the vignettes distributed with the package, which also provides greater detail on some of topics that we do describe, such as extensions to the package.

## 2. Preliminary Example: Polynomial regression

The data for the example in this section are drawn from the **ISLR2** package for R, associated with [James et al. \(2021\)](#). The presentation here is close (though not identical) to that in the original source ([James et al. 2021](#), Secs. 5.1, 5.3), and it demonstrates the use of the `cv()` function.<sup>1</sup>

The Auto data set contains information about 392 cars:

```
R> data("Auto", package="ISLR2")
R> summary(Auto)
```

mpg	cylinders	displacement	horsepower	weight
Min. : 9.0	Min. :3.00	Min. : 68	Min. : 46.0	Min. :1613
1st Qu.:17.0	1st Qu.:4.00	1st Qu.:105	1st Qu.: 75.0	1st Qu.:2225
Median :22.8	Median :4.00	Median :151	Median : 93.5	Median :2804
Mean :23.4	Mean :5.47	Mean :194	Mean :104.5	Mean :2978
3rd Qu.:29.0	3rd Qu.:8.00	3rd Qu.:276	3rd Qu.:126.0	3rd Qu.:3615
Max. :46.6	Max. :8.00	Max. :455	Max. :230.0	Max. :5140

acceleration	year	origin	name
Min. : 8.0	Min. :70	Min. :1.00	amc matador : 5
1st Qu.:13.8	1st Qu.:73	1st Qu.:1.00	ford pinto : 5
Median :15.5	Median :76	Median :1.00	toyota corolla : 5
Mean :15.5	Mean :76	Mean :1.58	amc gremlin : 4
3rd Qu.:17.0	3rd Qu.:79	3rd Qu.:2.00	amc hornet : 4
Max. :24.8	Max. :82	Max. :3.00	chevrolet chevette: 4
			(Other) :365

<sup>1</sup>James et al. (2021) use the `cv.glm()` function in the **boot** package (Canty and Ripley 2022; Davison and Hinkley 1997). Despite its name, `cv.glm()` is an independent function and not a method of a `cv()` generic function. The **boot** package is part of the standard R distribution.

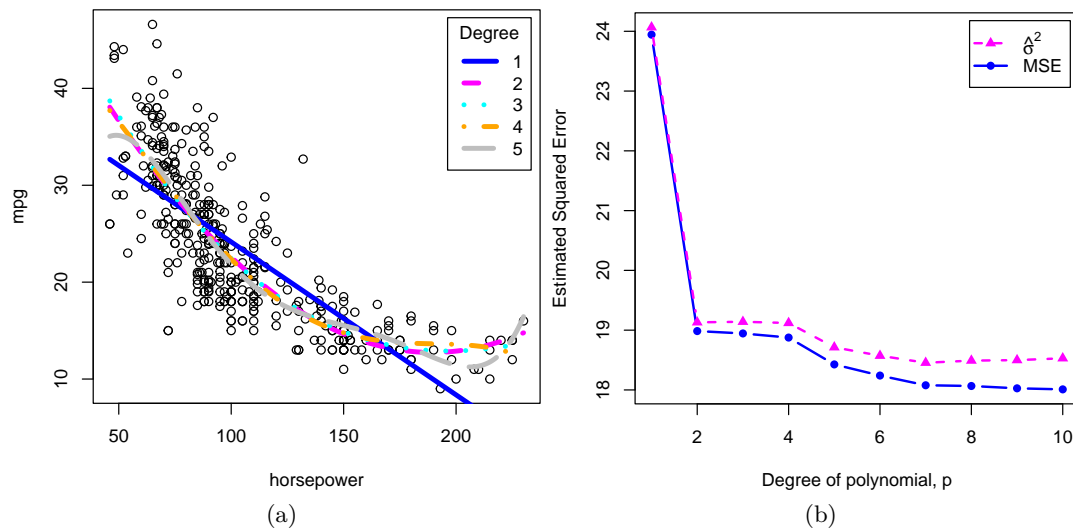


Figure 1: (a) `mpg` vs. `horsepower` for the `Auto` data, showing fitted polynomials of degree 1 through 5. (b) Estimated squared error as a function of polynomial degree,  $p = 1, \dots, 10$ .

With the exception of `origin` (which we don't use here), these variables are largely self-explanatory, except possibly for units of measurement: for details see `help("Auto", package="ISLR2")`.

We'll focus on the relationship of `mpg` (miles per gallon) to `horsepower`, as displayed in Figure 1 (a). The relationship between the two variables is monotone, decreasing, and nonlinear. Following James *et al.* (2021), we'll consider approximating the relationship by a polynomial regression, with the degree of the polynomial  $p$  ranging from 1 (a linear regression) to 10.<sup>2</sup> Polynomial fits for  $p = 1$  to 5 are shown in Figure 1 (a). The linear fit is clearly inappropriate; the fits for  $p = 2$  (quadratic) through 4 are very similar; and the fit for  $p = 5$  may over-fit the data by chasing one or two relatively high `mpg` values at the right (but see the CV results reported below).

Figure 1 (b) shows two measures of estimated (squared) error as a function of polynomial-regression degree: The mean-squared error ("MSE"), defined as  $\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$ , and the usual residual variance, defined as  $\hat{\sigma}^2 = \frac{1}{n-p-1} \sum_{i=1}^n (y_i - \hat{y}_i)^2$ . The former necessarily declines with  $p$  (or, more strictly, can't increase with  $p$ ), while the latter gets slightly larger for the largest values of  $p$ , with the "best" value, by a small margin, for  $p = 7$ .

The generic `cv()` function has an `"lm"` method,

```
R> args(cv::cv.lm)
```

```
function (model, data = insight::get_data(model), criterion = mse,
  k = 10L, reps = 1L, seed = NULL, details = k <= 10L, confint = n >=
    400L, level = 0.95, method = c("auto", "hatvalues", "Woodbury",
    "naive"), ncores = 1L, ...)
NULL
```

<sup>2</sup>Although it serves to illustrate the use of CV, a polynomial is not the best choice here. Consider, for example the scatterplot for log-transformed `mpg` and `horsepower`, produced by `plot(mpg ~ horsepower, data=Auto, log="xy")` (execution of which is left to the reader). We revisit the `Auto` data in Section 4.

which takes the following arguments:

- **model**, an "lm" object, the only required argument.
- **data**, which can usually be inferred from the **model** object.
- **criterion**, a function to compute the CV criterion (defaulting to **mse**).
- **k**, the number of folds to employ (defaulting to 10); the character value "n" or "loo" may be supplied to specify leave-one-out cross-validation.
- **reps**, the number of times to repeat the CV procedure (defaulting to 1).
- **seed**, the seed for R's pseudo-random number generator; if not specified a value is randomly selected, reported, and saved, so that the CV procedure is replicable.
- **confint**, whether or not to compute a confidence interval for the CV criterion, defaulting to TRUE if there are at least 400 cases; a confidence interval is computed only if the CV criterion can be expressed as the average of casewise components (see Section 7.2 for details).
- **level**, the level for the confidence interval (defaulting to 0.95).
- **method**, the computational method to employ: "hatvalues" is relevant only for LOO CV and bases computation on the hatvalues for the linear model; "Woodbury" employs the Woodbury matrix identity to compute the CV criterion with each fold deleted; "naive" updates the model using the **update()** function; and "auto" (the default) selects "hatvalues" for LOO CV and "Woodbury" for *k*-fold CV, both of which are much more efficient than literally updating the least-squares fit (see below and Section 7.1).
- **ncores**, the number of cores to employ for parallel computation; if **ncores** = 1 (the default), the computations are not parallelized.

To illustrate, we perform 10-fold CV for a quadratic polynomial fit to the Auto data:<sup>3</sup>

```
R> library("car", quietly=TRUE)
R>
R> m.auto <- lm(mpg ~ poly(horsepower, 2), data=Auto)
R> brief(m.auto)
```

	(Intercept)	poly(horsepower, 2)1	poly(horsepower, 2)2
Estimate	23.446	-120.14	44.09
Std. Error	0.221	4.37	4.37

Residual SD = 4.37 on 389 df, R-squared = 0.688

```
R> cv(m.auto, confint=TRUE)
```

---

<sup>3</sup>We load the **car** package (Fox and Weisberg 2019) here for the **brief()** function and to use below.

```
R RNG seed set to 692992
```

#### 10-Fold Cross Validation

```
method: Woodbury
criterion: mse
cross-validation criterion = 19.165
bias-adjusted cross-validation criterion = 19.156
95% CI for bias-adjusted CV criterion = (15.707, 22.604)
full-sample criterion = 18.985
```

The function reports the CV estimate of MSE, a bias-adjusted estimate of the MSE (the bias adjustment is explained in Section 7.2), and the MSE is also computed for the original, full-sample regression. Because the number of cases  $n = 392 < 400$  for the `Auto` data, we set the argument `confint=TRUE` to obtain a confidence interval for the MSE, which proves to be quite wide.

To perform LOO CV:

```
R> cv(m.auto, k="loo")
```

#### n-Fold Cross Validation

```
method: hatvalues
criterion: mse
cross-validation criterion = 19.248
```

The "hatvalues" method reports only the CV estimate of MSE. Alternative methods are to use the Woodbury matrix identity or the "naive" approach of literally refitting the model with each case omitted. All three methods produce exact results for a linear model (within the precision of floating-point computations); for example:

```
R> cv(m.auto, k="loo", method="naive", confint=TRUE)
```

#### n-Fold Cross Validation

```
method: naive
criterion: mse
cross-validation criterion = 19.248
bias-adjusted cross-validation criterion = 19.248
95% CI for bias-adjusted CV criterion = (15.779, 22.717)
full-sample criterion = 18.985
```

The "naive" and "Woodbury" methods also return the bias-adjusted estimate of MSE (and a confidence interval around it) along with the full-sample MSE, but bias isn't an issue for LOO CV.

This is a small regression problem and all three computational approaches are essentially instantaneous, but it is still of interest to investigate their relative speed. In the following comparison, we include the `cv.glm()` function from the **boot** package, which takes the naive approach, and for which we have to fit the linear model as an equivalent Gaussian GLM. We use the `microbenchmark()` function from the package of the same name ([Mersmann 2023](#)) for the timings:

```
R> m.auto.glm <- glm(mpg ~ poly(horsepower, 2), data=Auto)
R> boot::cv.glm(Auto, m.auto.glm)$delta # MSE, biased-corrected MSE
```

```
[1] 19.24821 19.24787
```

```
R> microbenchmark::microbenchmark(
+   hatvalues = cv(m.auto, k="loo"),
+   Woodbury = cv(m.auto, k="loo", method="Woodbury"),
+   naive = cv(m.auto, k="loo", method="naive"),
+   cv.glm = boot::cv.glm(Auto, m.auto.glm),
+   times=10
+ )
```

Warning in microbenchmark::microbenchmark(hatvalues = cv(m.auto, k = "loo"), : less accurate nanosecond times to avoid potential integer overflows

Unit: milliseconds

expr	min	lq	mean	median	uq	max
hatvalues	1.153863	1.185966	1.201644	1.200582	1.217126	1.255297
Woodbury	12.222223	12.384009	13.541341	12.596635	15.472949	16.394752
naive	213.958418	219.854874	243.479767	223.434297	272.095967	342.005723
cv.glm	378.032464	379.625970	394.416244	383.815329	389.642762	447.170723
neval cld						
10 a						
10 a						
10 b						
10 c						

On our computer, using the hatvalues is about an order of magnitude faster than employing Woodbury matrix updates, and more than two orders of magnitude faster than refitting the model.

## 2.1. Comparing competing models

The `cv()` function also has a method that can be applied to a list of regression models for the same data, composed using the `models()` function. For  $k$ -fold CV, the same folds are used for the competing models, which reduces random error in their comparison. This result can also be obtained by specifying a common seed for R's random-number generator while applying `cv()` separately to each model, but employing a list of models is more convenient for both  $k$ -fold and LOO CV (where there is no random component to the composition of the  $n$  folds).

We illustrate with the polynomial regression models of varying degree for the `Auto` data, beginning by fitting and saving the 10 models:

```
R> mlist <- vector(10, mode="list")
R> for (p in 1:10) mlist[[p]] <- lm(mpg ~ poly(horsepower, p), data = Auto)
R> names(mlist) <- paste0("m.", 1:10)
R> mlist[2] # e.g., the quadratic fit
```

\$m.2

Call:

```
lm(formula = mpg ~ poly(horsepower, p), data = Auto)
```

Coefficients:

```
      (Intercept)  poly(horsepower, p)1  poly(horsepower, p)2
              23.4                -120.1                  44.1
```

We then apply `cv()` to the list of 10 models (the `data` argument is required):

```
R> # 10-fold CV
R> mlist <- models(mlist) # create "modList" object
R>
R> cv.auto.10 <- cv(mlist, data=Auto, seed=2120)
R> cv.auto.10[2] # e.g., for quadratic model
```

Model m.2:

10-Fold Cross Validation

method: Woodbury

cross-validation criterion = 19.346

bias-adjusted cross-validation criterion = 19.327

full-sample criterion = 18.985

```
R> # LOO CV
R> cv.auto.loo <- cv(mlist, data=Auto, k="loo")
R> cv.auto.loo[2] # e.g., for quadratic model
```

Model m.2:

n-Fold Cross Validation

method: hatvalues

cross-validation criterion = 19.248

The `models()` function takes an arbitrary number of regression models as its arguments, which are optionally named, to create a "modList" object. Alternatively, as here, we can supply a pre-constructed, optionally named, list of models as the single argument to `models()`. To visualize the results, comparing the models, we invoke the `plot()` method for the "cvModList" object returned by `cv()` (see Figure 2):

```
R> plot(cv.auto.10, main="Polynomial Regressions, 10-Fold CV",
+       axis.args=list(labels=1:10), xlab="Degree of Polynomial, p")
R> plot(cv.auto.loo, main="Polynomial Regressions, LOO CV",
+       axis.args=list(labels=1:10), xlab="Degree of Polynomial, p")
```



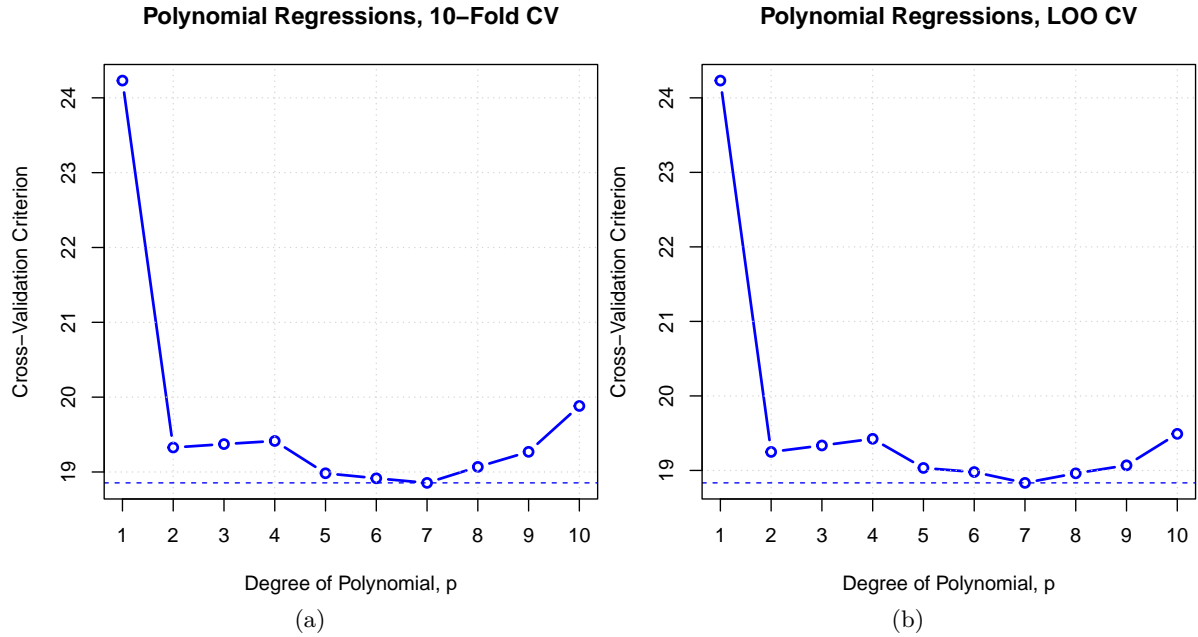


Figure 2: Cross-validated (a) 10-fold and (b) LOO MSE as a function of polynomial degree,  $p = 1, 2, \dots, 10$ .

In this example, 10-fold and LOO CV produce generally similar results, and also results that are similar to those produced by the estimated error variance  $\hat{\sigma}^2$  for each model (cf., Figure 1 (b) on page 4), except for the highest-degree polynomials, where the CV results more clearly suggest over-fitting.

### 3. Cross-validating mixed-effects models

The fundamental analogy for cross-validation is to the collection of new data. That is, predicting the response in each fold from the model fit to data in the other folds is like using the model fit to all of the data to predict the response for new cases from the values of the predictors for those new cases. As we explained, the application of this idea to independently sampled cases is straightforward.

In contrast, mixed-effects models are fit to *dependent* data, in which cases are clustered, such as hierarchical data, where the clusters comprise higher-level units (e.g., students clustered in schools), or longitudinal data, where the clusters are individuals and the cases are repeated observations on the individuals over time.<sup>4</sup>

<sup>4</sup>There are, however, more complex situations that give rise to so-called *crossed* (rather than *nested*) random effects. For example, consider students within classes within schools. In primary schools, students typically are in a single class, and so classes are nested within schools. In secondary schools, however, students typically take several classes and students who are together in a particular class may not be together in other classes; consequently, random effects based on classes within schools are crossed. The `lmer()` function in the `lme4` package, for example, is capable of modeling both nested and crossed random effects, and the `cv()` methods for mixed models in the `cv` package pertain to both nested and crossed random effects. We present an example of crossed random effects in a vignette for the `cv` package.

We can think of two approaches to applying cross-validation to clustered data:<sup>5</sup>

1. Treat CV as analogous to predicting the response for one or more cases in a *newly observed cluster*. In this instance, the folds comprise one or more whole clusters; we refit the model with all of the cases in clusters in the current fold removed; and then we predict the response for the cases in clusters in the current fold. These predictions are based only on fixed effects because the random effects for the omitted clusters are construed to be unknown, as they would be for data on cases in newly observed clusters.
2. Treat CV as analogous to predicting the response for a newly observed case in an *existing cluster*. In this instance, the folds comprise one or more individual cases, and the predictions can use both the fixed and random effects—so-called “best-linear-unbiased predictors” or “BLUPs.”

### 3.1. Example: The High-School and Beyond data

Following their use by [Raudenbush and Bryk \(2002\)](#), data from the 1982 *High School and Beyond* (HSB) survey have become a staple of the literature on mixed-effects models. The HSB data are used by [Fox and Weisberg \(2019, Sec. 7.2.2\)](#) to illustrate the application of linear mixed models to hierarchical data, and we’ll closely follow their example here.

The HSB data are included in the `MathAchieve` and `MathAchSchool` data sets in the `nlme` package ([Pinheiro and Bates 2000](#)). `MathAchieve` comprises individual-level data on 7185 students in 160 high schools, and `MathAchSchool` contains school-level data:

```
R> data("MathAchieve", package="nlme")
R> dim(MathAchieve)
```

```
[1] 7185    6
```

```
R> head(MathAchieve, 3)
```

```
Grouped Data: MathAch ~ SES | School
  School Minority    Sex    SES MathAch MEANSES
1   1224      No Female -1.528   5.876  -0.428
2   1224      No Female -0.588  19.708  -0.428
3   1224      No  Male -0.528  20.349  -0.428
```

```
R> tail(MathAchieve, 3)
```

```
Grouped Data: MathAch ~ SES | School
  School Minority    Sex    SES MathAch MEANSES
7183   9586      No Female  1.332  19.641   0.627
7184   9586      No Female -0.008  16.241   0.627
7185   9586      No Female  0.792  22.733   0.627
```

<sup>5</sup>We subsequently discovered that [Vehtari \(2023, Section 8\)](#) makes similar points.

```
R> data("MathAchSchool", package="nlme")
R> dim(MathAchSchool)
```

```
[1] 160    7
```

```
R> head(MathAchSchool, 2)
```

	School	Size	Sector	PRACAD	DISCLIM	HIMINTY	MEANSES
1224	1224	842	Public	0.35	1.597	0	-0.428
1288	1288	1855	Public	0.27	0.174	0	0.128

```
R> tail(MathAchSchool, 2)
```

	School	Size	Sector	PRACAD	DISCLIM	HIMINTY	MEANSES
9550	9550	1532	Public	0.45	0.791	0	0.059
9586	9586	262	Catholic	1.00	-2.416	0	0.627

The first few students are in school number 1224 and the last few in school 9586.

We'll use only the `School`, `SES` (students' socioeconomic status), and `MathAch` (their score on a standardized math-achievement test) variables in the `MathAchieve` data set, and `Sector` ("Catholic" or "Public") in the `MathAchSchool` data set.

Some data-management is required before fitting a mixed-effects model to the HSB data:

```
R> HSB <- MathAchieve
R> HSB <- merge(MathAchSchool[, c("School", "Sector")],
+             HSB[, c("School", "SES", "MathAch")], by="School")
R> names(HSB) <- tolower(names(HSB))
R> HSB <- within(HSB, {
+   mean.ses <- ave(ses, school)
+   cses <- ses - mean.ses
+ })
```

In the process, we merge variables from the school-level and student-level data sets, and create two new school-level variables: `mean.ses`, which is the average SES for students in each school; and `cses`, which is the individual students' SES centered at their school means. For details, see [Fox and Weisberg \(2019, Sec. 7.2.2\)](#).

Still following Fox and Weisberg, we proceed to use the `lmer()` function in the **lme4** package ([Bates, Mächler, Bolker, and Walker 2015](#)) to fit a mixed model for math achievement to the HSB data (summarizing the model with the `S()` function from the **car** package):

```
R> library("lme4", quietly=TRUE)
R> hsb.lmer <- lmer(mathach ~ mean.ses*cses + sector*cses
+                 + (cses | school), data=HSB)
R> S(hsb.lmer, brief=TRUE)
```

Estimates of Fixed Effects:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	12.128	0.199	60.86	< 2e-16
mean.ses	5.333	0.369	14.45	< 2e-16
cses	2.945	0.156	18.93	< 2e-16
sectorCatholic	1.227	0.306	4.00	6.2e-05
mean.ses:cses	1.039	0.299	3.48	0.00051
cses:sectorCatholic	-1.643	0.240	-6.85	7.3e-12

Estimates of Random Effects (Covariance Components):

Groups	Name	Std.Dev.	Corr
school	(Intercept)	1.543	
	cses	0.318	0.39
Residual		6.060	

Number of obs: 7185, groups: school, 160

logLik	df	AIC	BIC
-23252	10	46524	46592

We can then cross-validate at the cluster (i.e., school) level,

```
R> cv(hsb.lmer, k=10, clusterVariables="school", seed=5240)
```

R RNG seed set to 5240

10-Fold Cross Validation based on 160 {school} clusters

criterion: mse

cross-validation criterion = 39.15662

bias-adjusted cross-validation criterion = 39.14844

95% CI for bias-adjusted CV criterion = (38.06554, 40.23135)

full-sample criterion = 39.00599

or at the case (i.e., student) level,

```
R> cv(hsb.lmer, seed=1575)
```

R RNG seed set to 1575

Warning in checkConv(attr(opt, "derivs"), opt\$par, ctrl = control\$checkConv, :  
Model failed to converge with max|grad| = 0.00587207 (tol = 0.002, component 1)

boundary (singular) fit: see help('isSingular')

## 10-Fold Cross Validation

```

criterion: mse
cross-validation criterion = 37.44473
bias-adjusted cross-validation criterion = 37.33801
95% CI for bias-adjusted CV criterion = (36.28761, 38.38841)
full-sample criterion = 36.06767

```

For cluster-level CV, the `clusterVariables` argument tells `cv()` how the clusters are defined. Were there more than one clustering variable, say classes within schools, these would be provided as a character vector of variable names: `clusterVariables = c("school", "class")`. For cluster-level CV, the default is `k = "loo"`, that is, leave one cluster out at a time; we instead specify `k = 10` folds of clusters, each fold therefore comprising  $160/10 = 16$  schools.

If the `clusterVariables` argument is omitted, then case-level CV is employed, with `k = 10` folds as the default, here each with  $7185/10 \approx 719$  students. Notice that one of the 10 models refit with a fold removed failed to converge. Convergence problems are common in mixed-effects modeling. The issue here is that an estimated variance component is close to or equal to 0, which is at a boundary of the parameter space. That shouldn't disqualify the fitted model for the kind of prediction required for cross-validation.

`cv()` also has methods for mixed models fit by the `glmer()` function in the **lme4** package, the `lme()` function in the **nlme** package (Pinheiro and Bates 2000), and the `glmmTMB()` function in the **glmmTMB** package (Brooks, Kristensen, van Benthem, Magnusson, Berg, Nielsen, Skaug, Maechler, and Bolker 2017), along with a simple procedure for extending `cv()` to other classes of mixed-effects models. See the vignettes in the **cv** package for details.

### 3.2. Example: Contrasting cluster-based and case-based CV

We introduce four artificial data sets that exemplify aspects of cross-validation particular to hierarchical models. Using these data sets, we show that model comparisons employing cluster-based and those employing case-based cross-validation may not agree on a “best” model. Furthermore, commonly used measures of fit, such as mean-squared error, do not necessarily become smaller as models become larger, even when the models are nested, and even when the measure of fit is computed for the whole data set.

The four data sets differ in the magnitude of between-cluster variance compared to within-cluster variance. They serve to illustrate how fitting mixed models, and, consequently, the cross-validation of mixed models, is sensitive to relative variance, which affects the degree of shrinkage of within-cluster estimates of effects towards between-cluster estimates.<sup>6</sup>

Consider a researcher studying the effect of the dosage of a drug on the severity of symptoms

---

<sup>6</sup>The purpose of this section is to explain some of the complexities of cross-validating mixed-effects models, which motivates the approach that we take for mixed models in the **cv** package. The examples in this section introduce no features of the **cv** package that we haven't already encountered, and the generation of data for these examples involves unenlightening data management. Consequently, we don't show the code for data generation, model fitting, and graphs in this section; the code is in the complementary materials for the paper, however. For these analyses, we used the `glmmTMB()` function in the **glmmTMB** package (Brooks *et al.* 2017) because, in our experience, it is more likely to converge than functions in the **nlme** and the **lme4** packages for models with low between-cluster variance. We used the **lattice** (Sarkar 2008) and **latticeExtra** (Sarkar and Andrews 2022) packages for the graphs in this section.

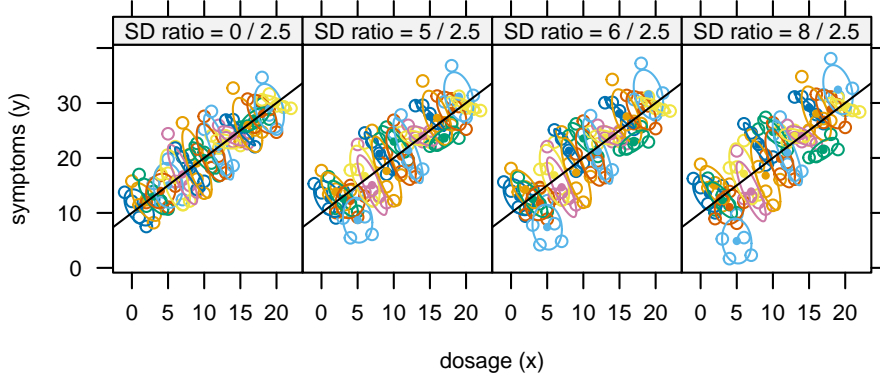


Figure 3: Data sets showing identical within-patient configurations with increasing between-patient variance. The labels above each panel show the between-patient SD/within-patient SD. The ellipses are estimated 50% concentration ellipses for each patient. The population between-patient regression line,  $E(y) = 10 + x$ , is shown in each panel, where  $y$  is symptoms and  $x$  is dosage.

for a hypothetical disease. The researcher has longitudinal data on 20 patients, each of whom was observed on 5 occasions in which patients received different dosages of the drug. The data are observational, with dosages prescribed by the patients' physicians, so that patients who were more severely affected by the disease received higher dosages of the drug.

Within patients, higher dosages are generally associated with a reduction in symptoms. Between patients, however, higher dosages are associated with higher levels of symptoms. A plausible mechanism is a reversal of causality: Within patients, higher dosages alleviate symptoms, but, between patients, higher morbidity causes the prescription of higher dosages.

To generate the four data sets, we first constructed a template data set from a population with a between-patient effect of dosage of 1.0 and a within-patient effect of  $-0.5$ . We then applied four different multipliers to the between-patient variation from the regression line of patient centroids. The resulting data sets, shown in Figure 3, differ in the between-patient variance of patient centroids from a common between-patient regression line, exhibiting a progression from low to high variance around the common regression line. Each patient has the same relative configuration of dosages and symptoms in each of the four data sets. To help visualize the structure of the data, estimated 50% Gaussian concentration ellipses are shown for each patient.

Using data like these, a researcher may attempt to obtain an estimate of the within-patient effect of dosage by fitting mixed models with a random intercept. We will illustrate how this approach results in estimates that are highly sensitive to the relative variance at the two levels of the mixed model by considering four models, all of which include a random intercept. The first three models are sequentially nested in their fixed effects: (1)  $\sim 1$ , intercept only; (2)  $\sim 1 + x$ , intercept and effect of  $x$ ; and (3)  $\sim 1 + x + x_m$ , intercept, effect of  $x$ , and a contextual variable,  $x_m$ , consisting of the within-patient mean of  $x$ . The fourth model,  $\sim 1 + I(x - x_m)$ , uses an intercept and the centered-within-group variable  $x - x_m$ . We thus fit four models to four data sets for a total of 16 models.

We proceed to obtain predictions from each model based on fixed effects alone, as would be used for cross-validation based on clusters (i.e., patients), and for fixed and random effects—

that is, the BLUPs—as would be used for cross-validation based on cases (i.e., occasions within patients):

The fixed- and random-effects predictions from these models are displayed in Figure 4 along with a summary of the data using 50% concentration ellipses (see the top row of the figure). Data sets with relatively low between-patient variance result in strong shrinkage of fixed-effects predictions, and also of BLUPs, towards the between-patient relationship between  $y$  and  $x$  in the model with an intercept and  $x$  as fixed-effect predictors,  $\sim 1 + x$ . The inclusion of a contextual variable in the model corrects this problem.

Although the BLUPs fit the observed data more closely than predictions based on fixed effects alone, the slopes of within-patient BLUPs do not conform to the within-patient slopes for the  $\sim 1 + x$  model in the two data sets with the smallest between-patient variances.

For data with a small between-patient variance, fixed-effects predictions for the  $\sim 1 + x$  model have a slope that is close to the between-patient slope but provide better overall predictions than the fixed-effect predictions for data sets with larger between-subject variance. With data whose between-patient variance is relatively large, predictions based on the model with a common intercept and slope for all clusters are very poor—indeed, much worse than the fixed-effects-only predictions based on the simpler random-intercept model.

We therefore anticipate (and show later in this section) that case-based cross-validation may prefer the intercept-only model,  $\sim 1$  to the larger  $\sim 1 + x$  model when the between-cluster variance is relatively small, but that cluster-based cross-validation will prefer the latter to the former.

We will discover that case-based cross-validation prefers the  $\sim 1 + x$  model to the  $\sim 1$  model for the ‘5 / 2.5’ data set, but cluster-based cross-validation prefers the latter model to the former. The situation is entirely reversed with the ‘8 / 2.5’ data set.

The third model,  $\sim 1 + x + xm$ , includes a contextual effect of  $x$ —that is, the cluster mean  $xm$ —along with  $x$  and the intercept in the fixed-effect part of the model, and with a random intercept. This model is equivalent to fitting  $y \sim I(x - xm) + xm + (1 | patient)$ , which is the model that generated the data. The fit of the mixed model  $\sim 1 + x + xm$  is consequently similar to that of a fixed-effects-only model with  $x$  and a categorical predictor for individual patients (i.e.,  $y \sim factor(patient) + x$ , treating patients as a factor, and not shown here).

We next carry out case-based cross-validation, which, as we have explained, is based on both fixed and predicted random effects (i.e., BLUPs), and cluster-based cross-validation, which is based on fixed effects only.<sup>7</sup> The results for all data sets, using both cluster-based and case-based cross-validation, are assembled and shown in Figure 5.

In summary, when between-cluster variance is relatively large, the model  $\sim 1 + x$ , with  $x$  alone and without the contextual mean of  $x$ , is assessed as fitting very poorly by cluster-based CV, but relatively much better by case-based CV. With all our examples, the model  $\sim 1 + x + xm$ , which includes both  $x$  and its contextual mean, produces better results using both cluster-based and case-based CV. These conclusions are consistent with our observations based on graphing predictions from the various models, and they illustrate the desirability of assessing mixed-effect models at different hierarchical levels.

---

<sup>7</sup>In order to reduce between-model random variability in comparisons of models on the same data set, we applied `cv()` to the list of models created by the `models()` function (introduced previously), performing cross-validation with the same folds for each model.

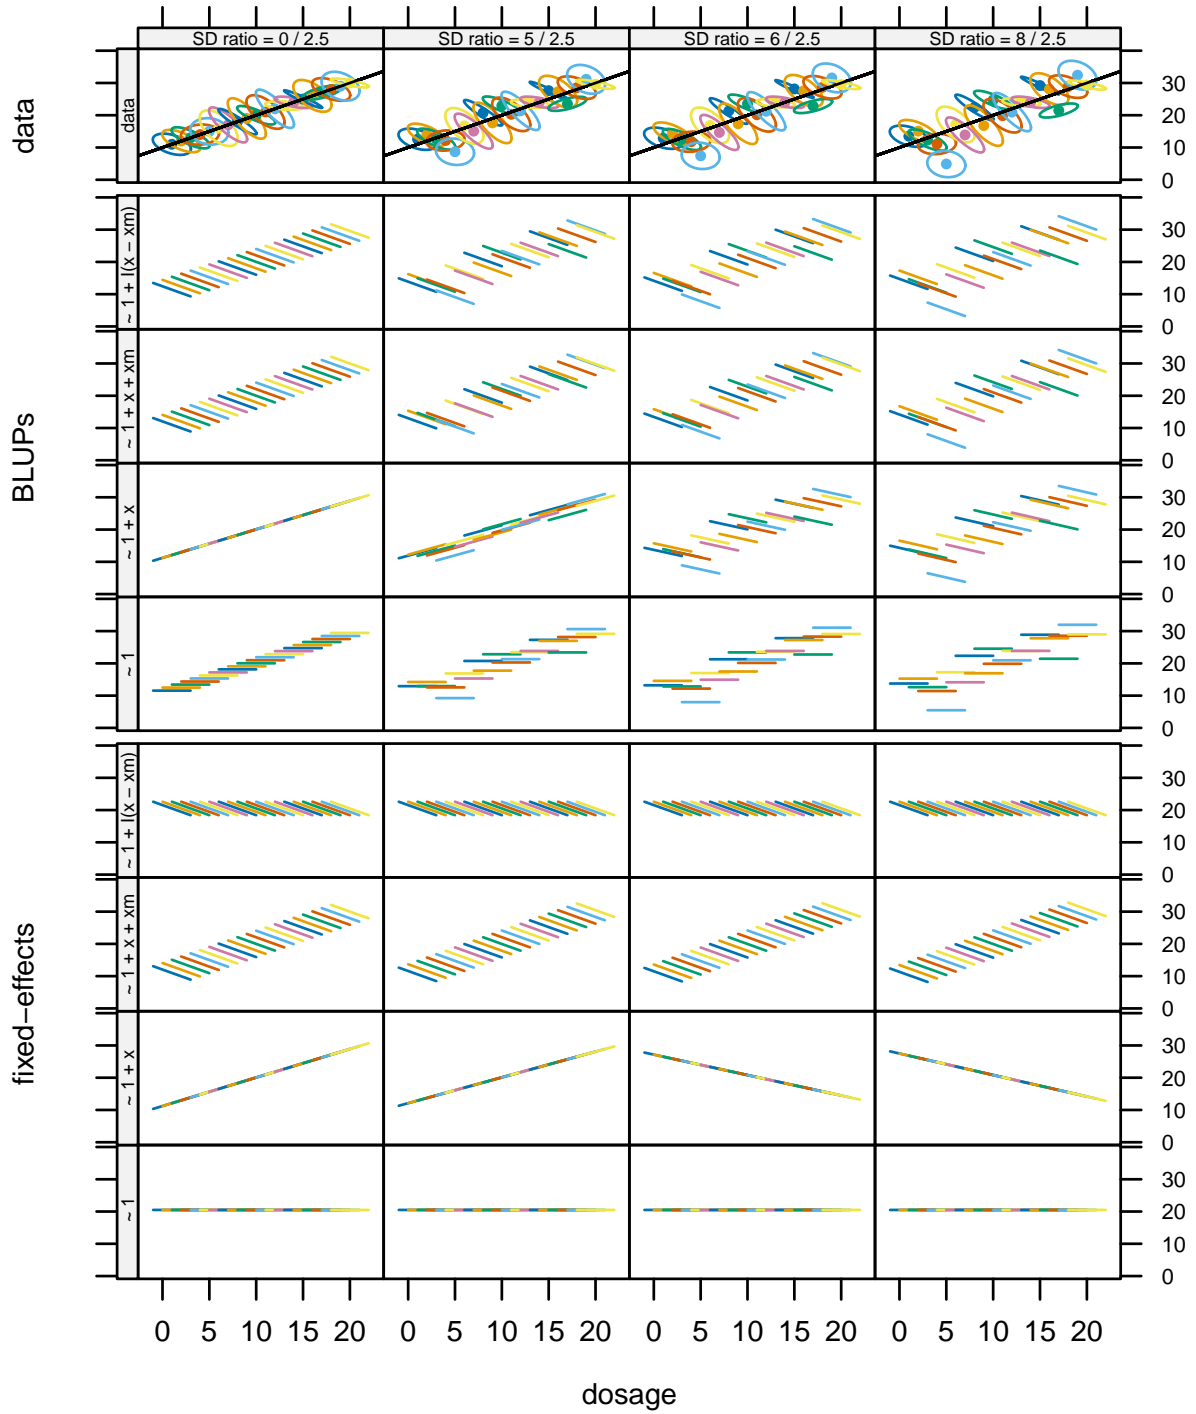


Figure 4: Fixed- and random-effect predictions (BLUPs) using each model applied to data sets with varying between- and within-patient variance ratios. The top row shows summaries of the within-patient data using estimated 50% concentration ellipses.






Figure 5: 10-fold cluster- and case-based cross-validation comparing mixed models with random intercepts and various fixed effects.

## 4. Cross-validating model specification

As [Hastie, Tibshirani, and Friedman \(2009, Sec. 7.10.2: “The Wrong and Right Way to Do Cross-validation”\)](#) explain, if the whole data are used to specify or fine-tune a statistical model, then subsequent cross-validation of the model is intrinsically misleading, because the model is selected to fit the whole data, including the part of the data that remains when each fold is removed. Statistical modeling is partly a craft, and one could imagine applying that craft to successive partial data sets, each with a fold removed. The resulting procedure would be tedious, though possibly worth the effort, but it would also be difficult to realize in practice: After all, we can hardly erase our memory of statistical modeling choices between analyzing partial data sets. Alternatively, if we’re able to automate the process of model specification, then we can more realistically apply CV mechanically.

The `function` method for `cv()` cross-validates a model-specification process in a general manner. Functions for four such model-specification processes are included in the package: `selectStepAIC()`, based on the `stepAIC()` function in the **MASS** package ([Venables and Ripley 2002](#)), performs stepwise predictor selection for regression models; `selectTrans()`, based on the `powerTransform()` function in the **car** package, transforms predictors and the response in a regression model towards normality; `selectTransAndStepAIC()`—the use of which we illustrate in the current section—performs both of these procedures sequentially; and `selectModelList()`—also illustrated in the current section—uses CV both to select one of several competing models, and then, recursively, to estimate prediction error for the selected model. In a vignette on extending the **cv** package, we explain how to specify additional model-selection procedures.

### 4.1. Example: Data transformation and predictor selection for the Auto data

To illustrate cross-validation of model specification, we return to the **Auto** data set:<sup>8</sup>

<sup>8</sup>This example benefits from an email conversation with Bill Venables, who of course isn’t responsible for

```

R> names(Auto)

[1] "mpg"          "cylinders"      "displacement"  "horsepower"    "weight"
[6] "acceleration" "year"          "origin"        "name"

R> xtabs(~ year, data=Auto)

year
70 71 72 73 74 75 76 77 78 79 80 81 82
29 27 28 40 26 30 34 28 36 29 27 28 30

R> xtabs(~ origin, data=Auto)

origin
  1   2   3
245 68 79

R> xtabs(~ cylinders, data=Auto)

cylinders
  3   4   5   6   8
  4 199   3  83 103

```

The `Auto` data appeared in a preliminary example in Section 2, where we employed CV to inform the selection of the degree of a polynomial regression of `mpg` on `horsepower`. Here, we consider more generally the problem of predicting `mpg` from the other variables in the `Auto` data. We begin with a bit of data management, and then examine the pairwise relationships among the numeric variables in the data set (Figure 6, produced by the `scatterplotMatrix()` function in the `car` package):

```

R> Auto$cylinders <- factor(Auto$cylinders,
+                           labels=c("3-4", "3-4", "5-6", "5-6", "8"))
R> Auto$year <- as.factor(Auto$year)
R> Auto$origin <- factor(Auto$origin,
+                        labels=c("America", "Europe", "Japan"))
R> rownames(Auto) <- make.names(Auto$name, unique=TRUE)
R> Auto$name <- NULL
R>
R> scatterplotMatrix(~ mpg + displacement + horsepower + weight + acceleration,
+                   smooth=list(spread=FALSE), data=Auto, pch=".")

```

A comment before we proceed: `origin` is clearly categorical and so converting it to a factor is natural, but we could imagine treating `cylinders` and `year` as numeric predictors. There are, however, only 5 distinct values of `cylinders` (ranging from 3 to 8), but cars with 3 or 5 cylinders are rare, and none of the cars has 7 cylinders. There are similarly only 13

---

the use to which we've put his insightful remarks.




Figure 6: Scatterplot matrix for the numeric variables in the `Auto` data. In each panel, the solid line shows the linear least-squares fit and the broken line is for a nonparametric regression.

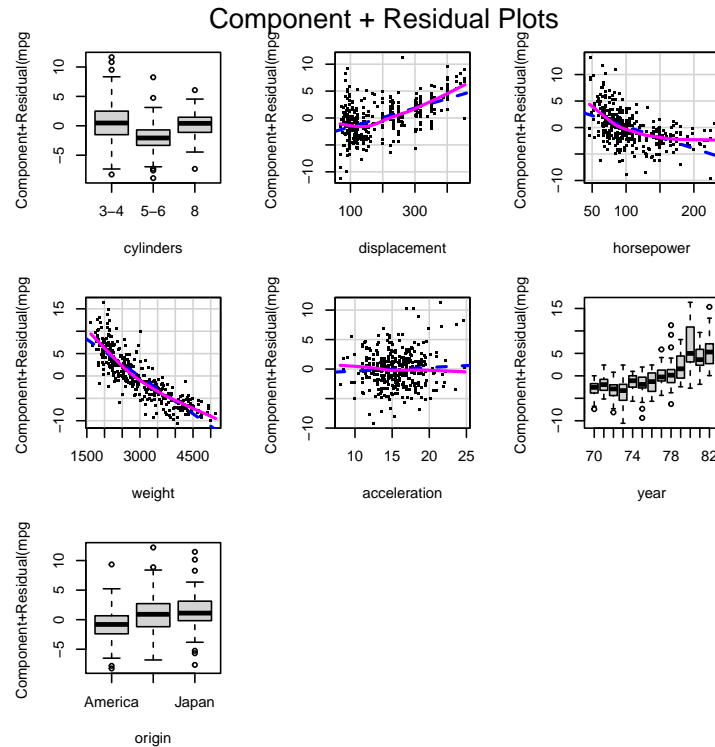


Figure 7: Component+residual plots for the working model fit to the **Auto** data. The broken blue lines show linear least-squares fits, while the solid magenta lines are nonparametric-regression smooths.

distinct years between 1970 and 1982 in the data, and the relationship between `mpg` and `year` is difficult to characterize.<sup>9</sup> It's apparent that most these variables are positively skewed and that many of the pairwise relationships among them are nonlinear.

We start with a “working model” that specifies linear partial relationships of the response to the numeric predictors:

```
R> m.auto <- lm(mpg ~ ., data = Auto)
R> crPlots(m.auto, pch=".")
```

The component+residual plots in Figure 7 clearly reveal the inadequacy of the model.

Some background: As [Weisberg \(2014, Sec. 8.2\)](#) explains, there are technical advantages to having (numeric) predictors in linear regression analysis that are themselves linearly related. If the predictors *aren't* linearly related, then the relationships between them can often be straightened by power transformations. Transformations can be selected after graphical examination of the data, or by analytic methods, such as transforming the predictors towards

<sup>9</sup>Making the decision to treat `year` as a factor on this basis could be construed as cheating in the current context, which illustrates the difficulty of automating the whole model-selection process. It's rarely desirable, in our opinion, to forgo exploration of the data to ensure the purity of model validation. We believe, however, that it's still useful to automate as much of the process as we can, as illustrated here, to obtain a more realistic, if still biased, estimate of the predictive power of a model.

multivariate normality, which implies linearity. Once the relationships between the predictors are linearized, it can be advantageous similarly to transform the conditional distribution of the response variable towards normality. Selecting transformations analytically raises the possibility of automating the process, as required for cross-validation.

The `powerTransform()` function in the `car` package transforms variables towards multivariate normality by a generalization of Box and Cox’s maximum-likelihood-like approach (Box and Cox 1964). Several “families” of power transformations can be used, including the original Box-Cox family (which is the default), simple powers (and roots), and two adaptations of the Box-Cox family to data that may include negative values and 0s: the Box-Cox-with-negatives family and the Yeo-Johnson family; see Weisberg (2014, Chap. 8) and Fox and Weisberg (2019, Chap. 3) for details. We proceed to transform the numeric predictors in the `Auto` regression towards multivariate normality:

```
R> num.predictors <- c("displacement", "horsepower", "weight",
+                      "acceleration")
R> tr.x <- powerTransform(Auto[, num.predictors])
R> summary(tr.x)
```

```
bcPower Transformations to Multinormality
              Est Power Rounded Pwr Wald Lwr Bnd Wald Up Bnd
displacement  -0.0509           0    -0.2082      0.1065
horsepower    -0.1249           0    -0.2693      0.0194
weight        -0.0870           0    -0.2948      0.1208
acceleration   0.3061           0    -0.0255      0.6376
```

```
Likelihood ratio test that transformation parameters are equal to 0
(all log transformations)
```

```
              LRT df  pval
LR test, lambda = (0 0 0 0) 4.8729  4 0.301
```

```
Likelihood ratio test that no transformations are needed
```

```
              LRT df  pval
LR test, lambda = (1 1 1 1) 390.08  4 <2e-16
```

We then apply the (rounded) transformations—all, as it turns out, logs (i.e., “Oth” powers)—to the data and re-estimate the model:

```
R> A <- Auto
R> powers <- tr.x$roundlam
R> for (pred in num.predictors){
+   A[, pred] <- bcPower(A[, pred], lambda=powers[pred])
+ }
R> m <- update(m.auto, data=A)
```

Having transformed the predictors towards multivariate normality, we now consider whether there’s evidence for transforming the response (using `powerTransform()` for Box and Cox’s original method), also obtaining a log transformation:

```
R> summary(powerTransform(m))
```

```
bcPower Transformation to Normality
```

```
  Est Power Rounded Pwr Wald Lwr Bnd Wald Up Bnd
Y1    0.0024          0   -0.1607    0.1654
```

```
Likelihood ratio test that transformation parameter is equal to 0
(log transformation)
```

```
              LRT df  pval
LR test, lambda = (0) 0.00080154  1 0.977
```

```
Likelihood ratio test that no transformation is needed
```

```
              LRT df  pval
LR test, lambda = (1) 124.13  1 <2e-16
```

```
R> m <- update(m, log(mpg) ~ .)
```

The transformed numeric variables are much better-behaved (cf., Figure 6, on page 19, and Figure 8):

```
R> scatterplotMatrix(~ log(mpg) + displacement + horsepower + weight
+                      + acceleration,
+                      smooth=list(spread=FALSE), data=A, pch=".")
```

And the partial relationships in the model fit to the transformed data are much more nearly linear (cf., Figure 7, on page 20, and Figure 9):

```
R> crPlots(m, pch=".")
```

After transforming both the numeric predictors and the response, we proceed to use the `stepAIC()` function in the **MASS** package to perform predictor selection, employing the BIC model-selection criterion (by setting the `k` argument of `stepAIC()` to  $\log n$ ):

```
R> library("MASS")
R> m.step <- stepAIC(m, k=log(nrow(A)), trace=FALSE)
R> summary(m.step)
```

Call:

```
lm(formula = log(mpg) ~ horsepower + weight + acceleration +
    year + origin, data = A)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-0.3523 -0.0568  0.0068  0.0674  0.3586
```




Figure 8: Scatterplot matrix for the transformed numeric variables in the Auto data




Figure 9: Component+residual plots for the model fit to the transformed Auto data

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	9.43459	0.26153	36.07	< 2e-16
horsepower	-0.27625	0.05614	-4.92	1.3e-06
weight	-0.60907	0.05600	-10.88	< 2e-16
acceleration	-0.13138	0.05319	-2.47	0.01397
year71	0.02798	0.02894	0.97	0.33412
year72	-0.00711	0.02845	-0.25	0.80274
year73	-0.03953	0.02601	-1.52	0.12947
year74	0.05275	0.02999	1.76	0.07936
year75	0.05320	0.02928	1.82	0.07004
year76	0.07432	0.02821	2.63	0.00878
year77	0.13793	0.02888	4.78	2.6e-06
year78	0.14588	0.02753	5.30	2.0e-07
year79	0.23604	0.02908	8.12	7.0e-15
year80	0.33527	0.03115	10.76	< 2e-16
year81	0.26287	0.03056	8.60	< 2e-16
year82	0.32339	0.02961	10.92	< 2e-16
originEurope	0.05582	0.01678	3.33	0.00097
originJapan	0.04355	0.01748	2.49	0.01314

Residual standard error: 0.105 on 374 degrees of freedom

Multiple R-squared: 0.909, Adjusted R-squared: 0.905

F-statistic: 220 on 17 and 374 DF, p-value: <2e-16

The selected model includes three of the numeric predictors, `horsepower`, `weight`, and `acceleration`, along with the factors `year` and `origin`. We can calculate the MSE for this model, but we expect that the result will be optimistic because we used the whole data to help specify the model:

```
R> mse(Auto$mpg, exp(fitted(m.step)))
```

```
[1] 6.5121
attr("casewise loss")
[1] "(y - yhat)^2"
```

This is considerably smaller than the MSE for the original working model:

```
R> mse(Auto$mpg, fitted(m.auto))
```

```
[1] 8.0932
attr("casewise loss")
[1] "(y - yhat)^2"
```

A perhaps subtle point is that we compute the MSE for the selected model on the original `mpg` response scale rather than the log scale, so as to make the selected model comparable to



the working model.<sup>10</sup>

The "function" method for `cv()` allows us to cross-validate the whole model-selection procedure. The first argument to `cv()`—here `selectTransAndStepAIC`—is a model-selection function capable of refitting the model with a fold omitted and returning a CV criterion:

```
R> num.predictors
```

```
[1] "displacement" "horsepower"    "weight"          "acceleration"
```

```
R> cvs <- cv(selectTransStepAIC, data=Auto, seed=76692,
+           working.model=m.auto,
+           predictors=num.predictors,
+           response="mpg", AIC=FALSE)
```

```
R RNG seed set to 76692
```

```
R> cvs
```

```
10-Fold Cross Validation
```

```
criterion: mse
```

```
cross-validation criterion = 7.4856
```

```
bias-adjusted cross-validation criterion = 7.3435
```

```
full-sample criterion = 6.5121
```

```
R> compareFolds(cvs)
```

```
CV criterion by folds:
```

fold.1	fold.2	fold.3	fold.4	fold.5	fold.6	fold.7	fold.8	fold.9	fold.10
6.0006	10.5093	7.4444	10.6179	5.0492	7.1832	4.0777	10.5019	9.1812	4.2507

```
Coefficients by folds:
```

	(Intercept)	horsepower	lam.acceleration	lam.displacement	lam.horsepower
Fold 1	9.71384	-0.17408	0.50000	0.00000	0.00000
Fold 2	9.21713	-0.31480	0.00000	0.00000	0.00000
Fold 3	9.61824	-0.19248	0.00000	0.00000	0.00000
Fold 4	8.69910	-0.25523	0.50000	0.00000	0.00000
Fold 5	9.14403	-0.14934	0.00000	0.00000	0.00000
Fold 6	9.63481	-0.16739	0.50000	0.00000	0.00000
Fold 7	9.98933	-0.36847	0.00000	0.00000	-0.15447
Fold 8	9.06301	-0.29721	0.00000	0.00000	0.00000
Fold 9	8.88315	-0.22684	0.00000	0.00000	0.00000
Fold 10	9.61727	-0.17086	0.00000	0.00000	0.00000

<sup>10</sup>That's slightly uncomfortable given the skewed distribution of `mpg`. An alternative is to use a robust measure of model lack-of-fit, such as the median absolute error instead of the mean-squared error, employing the `medAbsErr()` function from the `cv` package. The median absolute error, however, cannot be expressed as a casewise average (see Section 7.2).

	lam.weight	lambda	weight	year71	year72	year73	year74	
Fold 1	0.00000	0.00000	-0.74636	0.03764	-0.00327	-0.02477	0.05606	
Fold 2	0.00000	0.00000	-0.47728	0.02173	-0.01488	-0.03770	0.04312	
Fold 3	0.00000	0.00000	-0.72085	0.01128	-0.02569	-0.03872	0.05187	
Fold 4	0.00000	0.00000	-0.53846	0.02153	-0.02922	-0.05181	0.04136	
Fold 5	0.00000	0.00000	-0.69081	0.02531	-0.01062	-0.04625	0.05039	
Fold 6	0.00000	0.00000	-0.74049	0.02456	0.00759	-0.03412	0.06266	
Fold 7	0.00000	0.00000	-0.72843	0.02532	-0.01271	-0.04144	0.04568	
Fold 8	0.00000	0.00000	-0.46392	0.02702	-0.02041	-0.05605	0.04437	
Fold 9	0.00000	0.00000	-0.47136	0.00860	-0.03620	-0.04835	0.01906	
Fold 10	0.00000	0.00000	-0.73550	0.02937	-0.00899	-0.03814	0.05408	
	year75	year76	year77	year78	year79	year80	year81	year82
Fold 1	0.07080	0.07250	0.14420	0.14281	0.23266	0.35127	0.25635	0.30546
Fold 2	0.04031	0.06718	0.13094	0.14917	0.21871	0.33192	0.26196	0.30943
Fold 3	0.03837	0.06399	0.11593	0.12601	0.20499	0.32821	0.24478	0.29204
Fold 4	0.04072	0.05537	0.12292	0.14083	0.22878	0.32947	0.25140	0.27248
Fold 5	0.05596	0.07044	0.13356	0.14724	0.24675	0.33331	0.26938	0.32594
Fold 6	0.06940	0.07769	0.14211	0.14647	0.23532	0.34761	0.26737	0.33062
Fold 7	0.03614	0.07385	0.12976	0.14040	0.23976	0.33998	0.27652	0.30659
Fold 8	0.06573	0.08135	0.13158	0.13987	0.23011	0.32880	0.25886	0.30538
Fold 9	0.03018	0.05846	0.10536	0.11722	0.20665	0.31533	0.23352	0.29375
Fold 10	0.04881	0.07862	0.14101	0.14313	0.23258	0.35649	0.26214	0.32421
	acceleration	displacement	cylinders5-6	cylinders8	originEurope			
Fold 1								
Fold 2	-0.18909		-0.09197					
Fold 3								
Fold 4	-0.03484			-0.09080	-0.10909			
Fold 5						0.06261		
Fold 6								
Fold 7								
Fold 8	-0.17676		-0.10542					
Fold 9	-0.14514		-0.13452					
Fold 10								
	originJapan							
Fold 1								
Fold 2								
Fold 3								
Fold 4								
Fold 5	0.04							
Fold 6								
Fold 7								
Fold 8								
Fold 9								
Fold 10								

The other arguments to `cv()` are (see `?cv:::cv.function` for additional optional arguments and details):

- **data**, the data set to which the model is fit.
- **seed**, an optional seed for R’s pseudo-random-number generator; as for `cv()`, if the seed isn’t supplied by the user, then a seed is randomly selected and saved.
- Arguments required by the model-selection function: the starting **working.model** (here, **m.auto**) for transformation and predictor selection; the names of the variables—**predictors** and **response**—that are candidates for transformation; and **AIC=FALSE**, which specifies use of the BIC for model selection.

Some noteworthy points:

- `selectTransStepAIC()` automatically computes CV cost criteria, here the MSE, on the *untransformed* response scale.
- As we anticipated, the estimate of the MSE that we obtain by cross-validating the whole model-specification process is larger than the MSE computed for the model we fit to the Auto data separately selecting transformations of the predictors and the response and then selecting predictors for the whole data set.
- When we look at the transformations and predictors selected with each of the 10 folds omitted (i.e., the output of `compareFolds()`), we see that there is little uncertainty in choosing variable transformations (the `lam.*s` for the *xs* and `lambda` for *y* in the output), but considerably more uncertainty in subsequently selecting predictors: **horsepower**, **weight**, and **year** are always included among the selected predictors; **acceleration** and **displacement** are included respectively in 4 and 3 of 10 selected models; and **cylinders** and **origin** are each included in only 1 of 10 models. Recall that when we selected predictors for the full data, we obtained a model with **horsepower**, **weight**, **acceleration**, **year**, and **origin**.

## 4.2. Example: Applying recursive CV to polynomial regression for the Auto data<sup>11</sup>

In Section 2.1, following James *et al.* (2021, Secs. 5.1, 5.3), we fit polynomial regressions up to degree 10 to the relationship of **mpg** to **horsepower** for the Auto data, saving the 10 models, named **m.1** through **m.10**, in the list **m.list**. We then used `cv()` to compare the cross-validated MSE for the 10 models, discovering that the 7th degree polynomial had the smallest MSE (by a small margin).

If we select the 7th degree polynomial model, intending to use it for prediction, the CV estimate of the MSE for this model will be optimistic. One solution is to cross-validate the process of using CV to select the “best” model—that is, to apply CV to CV recursively. The function `selectModelList()`, which is suitable for use with `cv()`, implements this idea.

Applying `selectModelList()` to the Auto polynomial-regression models, and using 10-fold CV, we obtain:

```
R> recursiveCV.auto <- cv(selectModelList, data=Auto,
+                          working.model=m.list,
+                          save.model=TRUE, seed=2120)
```

<sup>11</sup>What we call here “recursive CV” has also been termed “nested CV” (e.g., in the **cvms** package by Olsen and Zachariae 2024). We prefer, however, to reserve the term “nested CV” for the procedure described by Bates, Hastie, and Tibshirani (2023).

```
R RNG seed set to 2120
```

```
R> recursiveCV.auto
```

```
10-Fold Cross Validation
```

```
criterion: mse
```

```
cross-validation criterion = 20.012
```

```
bias-adjusted cross-validation criterion = 20.619
```

```
full-sample criterion = 18.746
```

```
R> recursiveCV.auto$selected.model
```

```
Call:
```

```
lm(formula = mpg ~ poly(horsepower, p), data = Auto)
```

```
Coefficients:
```

(Intercept)	poly(horsepower, p)1	poly(horsepower, p)2
23.45	-120.14	44.09
poly(horsepower, p)3	poly(horsepower, p)4	poly(horsepower, p)5
-3.95	-5.19	13.27
poly(horsepower, p)6	poly(horsepower, p)7	
-8.55	7.98	

```
R> cv(mlist[[7]], seed=2120) # same seed for same folds
```

```
R RNG seed set to 2120
```

```
10-Fold Cross Validation
```

```
method: Woodbury
```

```
criterion: mse
```

```
cross-validation criterion = 18.898
```

```
bias-adjusted cross-validation criterion = 18.854
```

```
full-sample criterion = 18.078
```

As expected, recursive CV produces a larger estimate of MSE for the selected 7th degree polynomial model than CV applied directly to this model.

We can equivalently call `cv()` with the list of models as its first argument and set the argument `recursive=TRUE`:

```
R> cv(mlist, data=Auto, seed=2120, recursive=TRUE, save.model=TRUE)
```

```
R RNG seed set to 2120
```

```
10-Fold Cross Validation
```

```
cross-validation criterion = 20.012
```

```
bias-adjusted cross-validation criterion = 20.619
```

```
full-sample criterion = 18.746
```

## 5. Extending the `cv` package

The `cv` package is designed to be extensible in several directions. In order of increasing general complexity, we can add: (1) a cross-validation cost criterion; (2) a model class that's not directly accommodated by the `cv()` default method or by another directly inherited method; and (3) a new model-selection procedure suitable for use with the `"function"` method for `cv()`. In this section, we illustrate (1) and (2); more diverse and extensive examples may be found in the vignette on extending the `cv` package.

Suppose that we want to cross-validate a multinomial logistic regression model fit by the `multinom()` function in the `nnet` package (Venables and Ripley 2002). We borrow an example from Fox (2016, Sec. 14.2.1), with data from the British Election Panel Study on vote choice in the 2001 British election. Data for the example are in the `BEPS` data set in the `carData` package:

```
R> data("BEPS", package="carData")
R> summary(BEPS)
```

	vote	age	economic.cond.national	
Conservative	:462	Min. :24.0	Min. :	1.00
Labour	:720	1st Qu.:41.0	1st Qu.:	3.00
Liberal Democrat	:343	Median :53.0	Median :	3.00
		Mean :54.2	Mean :	3.25
		3rd Qu.:67.0	3rd Qu.:	4.00
		Max. :93.0	Max. :	5.00
	economic.cond.household	Blair	Hague	Kennedy
Min. :	1.00	Min. :	1.00	Min. :1.00
1st Qu.:	3.00	1st Qu.:	2.00	1st Qu.:2.00
Median :	3.00	Median :	4.00	Median :3.00
Mean :	3.14	Mean :	3.33	Mean :2.75
3rd Qu.:	4.00	3rd Qu.:	4.00	3rd Qu.:4.00
Max. :	5.00	Max. :	5.00	Max. :5.00
	Europe	political.knowledge	gender	
Min. :	1.00	Min. :	0.00	female:812
1st Qu.:	4.00	1st Qu.:	0.00	male :713
Median :	6.00	Median :	2.00	
Mean :	6.73	Mean :	1.54	
3rd Qu.:	10.00	3rd Qu.:	2.00	
Max. :	11.00	Max. :	3.00	

The polytomous (multi-category) response variable is `vote`, a factor with levels `"Conservative"`, `"Labour"`, and `"Liberal Democrat"`. The predictors of `vote` are:

- `age`, in years;
- `econ.cond.national` and `econ.cond.household`, the respondent's ratings of the state of the economy, on 1 to 5 scales.
- `Blair`, `Hague`, and `Kennedy`, ratings of the leaders of the Labour, Conservative, and Liberal Democratic parties, on 1 to 5 scales.

- **Europe**, an 11-point scale on attitude towards European integration, with high scores representing “Euro-skepticism.”
- **political.knowledge**, knowledge of the parties’ positions on European integration, with scores from 0 to 3.
- **gender**, “female” or “male”.

The model fit to the data includes an interaction between **Europe** and **political.knowledge**, which was the focus of the original research on which this example is based (Andersen, Heath, and Sinnott 2002); the other predictors enter the model additively:

```
R> library("nnet", quietly=TRUE)
R> m.beps <- multinom(vote ~ age + gender + economic.cond.national +
+                      economic.cond.household + Blair + Hague + Kennedy +
+                      Europe*political.knowledge, data=BEPS)

# weights:  36 (22 variable)
initial  value 1675.383740
iter   10 value 1240.047788
iter   20 value 1163.199642
iter   30 value 1116.519687
final   value 1116.519666
converged
```

The **Europe**  $\times$  **political.knowledge** interaction is associated with a very small *p*-value. Figure 10 shows an “effect plot,” using the **effects** package (Fox and Weisberg 2019) to visualize the interaction in a “stacked-area” graph:

```
R> plot(effects::Effect(c("Europe", "political.knowledge"), m.beps,
+                      xlevels=list(Europe=1:11, political.knowledge=0:3),
+                      fixed.predictors=list(given.values=c(gendermale=0.5))),
+       lines=list(col=c("blue", "red", "orange")), # party colors
+       axes=list(x=list(rug=FALSE), y=list(style="stacked")))
```

As political knowledge increases, voters tend to align their votes more closely with the party positions on European integration: The Conservative Party was relatively Euro-skeptic, while Labour and the Liberal Democrats were more supportive of the UK’s participation in the EU. To cross-validate this multinomial-logit model we need an appropriate cost criterion. None of the criteria in the **cv** package will do—for example, **mse()** is appropriate only for a numeric response. The **BayesRule()** criterion, also supplied by **cv**, which is for a binary response, comes close:

```
R> BayesRule

function(y, yhat) {
  if (!all(y %in% c(0, 1)))
    stop("response values not all 0 or 1")
}
```




Figure 10: Effect plot for the interaction between attitude towards European integration and political knowledge in the multinomial logit model fit to voting data from the 2001 British Election Panel Study.

```

if (any(yhat < 0) ||
    any(yhat > 1))
  stop("fitted values outside of interval [0, 1]")
yhat <- round(yhat)
result <- mean(y != yhat) # proportion in error
attr(result, "casewise loss") <- "y != round(yhat)"
result
}
<bytecode: 0x1341aeac8>
<environment: namespace:cv>

```

After doing some error checking, `BayesRule()` rounds the predicted probability of a 1 (“success”) response in a binary regression model to 0 or 1 to obtain a categorical prediction, and then reports the proportion of incorrect predictions. Because the Bayes’s rule criterion is an average of casewise components (as, e.g., is the MSE), a “casewise loss” attribute is attached to the result, enabling the computation of bias correction and confidence intervals (as discussed in Section 7.2).

It is straightforward to adapt Bayes’s rule to a polytomous response:

```

R> head(BEPS$vote)

[1] Liberal Democrat Labour          Labour          Labour
[5] Labour          Labour
Levels: Conservative Labour Liberal Democrat

R> yhat <- predict(m.beps, type="class")
R> head(yhat)

```

```
[1] Labour          Labour          Labour          Labour
[5] Liberal Democrat Labour
Levels: Conservative Labour Liberal Democrat
```

```
R> BayesRuleMulti <- function(y, yhat){
+   result <- mean(y != yhat)
+   attr(result, "casewise loss") <- "y != yhat"
+   result
+ }
R>
R> BayesRuleMulti(BEPS$vote, yhat)

[1] 0.31869
attr("casewise loss")
[1] "y != yhat"
```

The `predict()` method for "multinom" models called with argument `type="class"` reports the Bayes's rule prediction for each case—that is, the response category with the highest predicted probability. Our `BayesRuleMulti()` function calculates the proportion of misclassified cases. Because this value is also the mean of casewise components, we attach a "casewise loss" attribute to the result.

The marginal proportions for the response categories are

```
R> xtabs(~ vote, data=BEPS)/nrow(BEPS)

vote
      Conservative      Labour Liberal Democrat
          0.30295          0.47213          0.22492
```

and so the marginal Bayes's rule prediction, that everyone will vote Labour, produces an error rate of  $1 - 0.47213 = 0.52787$ . The multinomial-logit model appears to do substantially better than that, but does its performance hold up to cross-validation?

We check first whether the default `cv()` method works “out-of-the-box” for the "multinom" model:

```
R> cv(m.beps, seed=3465, criterion=BayesRuleMulti)
```

```
Error in GetResponse.default(model): non-vector response
```

The default method of `GetResponse()` (a function supplied by the `cv` package—see `?GetResponse`) fails for a "multinom" object. A straightforward solution is to supply a `GetResponse.multinom()` method that returns the factor response (using the `get_response()` function from the **insight** package, Lüdtke, Waggoner, and Makowski 2019),

```
R> GetResponse.multinom <- function(model, ...) {
+   insight::get_response(model)
+ }
R>
R> head(GetResponse(m.beps))
```



```
[1] Liberal Democrat Labour          Labour          Labour
[5] Labour          Labour
Levels: Conservative Labour Liberal Democrat
```

and to try again:

```
R> cv(m.beps, seed=3465, criterion=BayesRuleMulti)
```

```
R RNG seed set to 3465
```

```
# weights: 36 (22 variable)
initial value 1507.296060
iter 10 value 1134.575036
iter 20 value 1037.413231
iter 30 value 1007.705242
iter 30 value 1007.705235
iter 30 value 1007.705235
final value 1007.705235
converged
```

```
Error in match.arg(type): 'arg' should be one of "class", "probs"
```

A `traceback()` (not shown) reveals that the problem is that the default method of `cv()` calls the "multinom" method for `predict()` with the argument `type="response"`, when the correct argument should be `type="class"`. We therefore must write a "multinom" method for `cv()`, but that proves to be very simple:

```
R> cv.multinom <- function (model, data, criterion = BayesRuleMulti,
+                           k, reps, seed, ...) {
+   model <- update(model, trace = FALSE)
+   NextMethod(
+     type = "class",
+     criterion = criterion,
+     criterion.name = deparse(substitute(criterion))
+   )
+ }
```

That is, we simply call the default `cv()` method (via `NextMethod()`) with the `type` argument properly set. In addition to supplying the correct `type` argument, our method sets the default `criterion` for the `cv.multinom()` method to `BayesRuleMulti`. Adding the argument `criterion.name = deparse(substitute(criterion))` is inessential, but it insures that printed output will include the name of the criterion function that's employed, whether it's the default `BayesRuleMulti` or something else. Prior to invoking `NextMethod()`, we called `update()` with `trace=FALSE` to suppress the iteration history reported by default by `multinom()`—it would be tedious to see the iteration history for each fold.

Then:

```
R> cv(m.beps, seed=3465)
```

```
R RNG seed set to 3465
```

```
10-Fold Cross Validation
```

```
criterion: BayesRuleMulti
```

```
cross-validation criterion = 0.32459
```

```
bias-adjusted cross-validation criterion = 0.32368
```

```
95% CI for bias-adjusted CV criterion = (0.30017, 0.34718)
```

```
full-sample criterion = 0.31869
```

The cross-validated polytomous Bayes’s rule criterion confirms that the fitted model does substantially better than the marginal Bayes’s rule prediction that everyone votes for Labour.

## 6. Comparing **cv** to other R software for cross-validation

The **cv** package is far from unique in implementing cross-validation of regression models in R. We’ve already mentioned the `cv.glm()` function in the **boot** package. A general review of CV facilities in R is beyond the scope of the current paper, but in this section we remark selectively on other R packages that support CV.

Our goal in implementing CV is to make it conveniently available to R users with limited programming skills,<sup>12</sup> and hence to encourage its use. Towards this end, the **cv** package provides a simple, consistent interface through the `cv()` generic function, with specific methods for different classes of standard R statistical models. Moreover, we have tried to make writing methods for additional classes of statistical models as simple and straightforward as possible. Because CV is computationally intensive, we also aim for efficiency, by enabling parallel computations generally, and by exploiting properties of specific classes of statistical models (in particular, computations based on hatvalues and the Woodbury matrix identity for linear and generalized linear model). Finally, the package includes some unusual features, such as general support for mixed-effects models and for cross-validating complex model-specification procedures.

Some R packages for statistical learning include facilities for cross-validation. Notable examples are the **caret** package (Kuhn 2008) and the **tidyfit** package (Pfzinger 2024); the latter employs **rsample** (Frick, Chow, Kuhn, Mahoney, Silge, and Wickham 2024) for CV, which we will discuss presently. This contrasts with the approach taken in the **cv** package, which, as we have explained, directly supports classes of standard R statistical models.

Other packages, including **cvms** (Olsen and Zachariae 2024), **mlexperiments** (Kapsner 2024), **origami** (Coyle, Hejazi, Malenica, and Phillips 2022), and **rsample** (which also supports bootstrapping), modularize the CV process to a greater or lesser extent. Advantages of this approach include flexibility and generality, but a disadvantage is that use of these packages entails nontrivial programming effort and skill.

We illustrate with an example adapted from a vignette in the **rsample** package, which uses the **attrition** data set from the **modeldata** package (Kuhn 2024), comparing LOO CV computed

---

<sup>12</sup>R users with sophisticated programming skills would generally find it unchallenging to implement CV directly for specific applications (as we illustrate below). Even in this case, however, there’s an argument for using a simple pre-programmed interface to CV, to minimize programming effort and to avoid mistakes.

using the `cv()` function in the **cv** package to LOO CV computed using the **caret** and **rsample** packages. At more than 15 million cumulative downloads, **caret** is by a very wide margin the most downloaded from the RStudio CRAN mirror of the R packages mentioned above. Among the packages that focus more on cross-validation, **rsample** is the most downloaded, at nearly 4 million downloads.

Here, with minimal explanation, is how one would perform LOO CV for this example using **rsample**.<sup>13</sup>

```
R> library("rsample")
R>
R> data("attrition", package = "modeldata")
R> nrow(attrition)

[1] 1470

R> # formula for model to be fit; Attrition is binary:
R> mod_form <- Attrition ~ JobSatisfaction + Gender + MonthlyIncome
R>
R> rs_obj <- loo_cv(attrition) # original uses vfold_cv()
R>
R> # user-supplied function:
R> holdout_results <- function(splits, ...) {
+   mod <- glm(..., data=analysis(splits), family=binomial)
+   holdout <- assessment(splits)
+   res <- broom::augment(mod, newdata=holdout)
+   lvls <- levels(holdout$Attrition)
+   predictions <- factor(ifelse(res$.fitted > 0, lvls[2], lvls[1]),
+                           levels=lvls)
+   res$correct <- predictions == holdout$Attrition
+   res
+ }
R>
R> library("purrr", warn.conflicts=FALSE)
R> # apply CV:
R> rs_obj$results <- map(rs_obj$splits, holdout_results, mod_form)
R> # summarize results:
R> rs_obj$accuracy <- map_dbl(rs_obj$results,
+                             function(x) mean(x$correct))
R> # accuracy is the complement of the Bayes-rule CV criterion:
R> 1 - summary(rs_obj$accuracy)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	1.0000	0.0000	0.0000	0.1612	0.0000	0.0000

<sup>13</sup>The example in the **rsample** vignette uses 10-fold rather than LOO CV, but otherwise the code from the vignette is only trivially modified here.

In this case, it's simple (and, in our opinion, simpler) to program LOO CV directly. For example

```
R> mod.attrition <- glm(mod_form, data=attrition, family=binomial)
R> n <- nrow(attrition)
R> yhat <- numeric(n)
R> for (i in 1:n){
+   m <- update(mod.attrition, data=attrition[-i, ])
+   yhat[i] <- predict(m, newdata=attrition[i, ], type="response")
+ }
R> mean(((attrition$Attrition == "Yes") - round(yhat))^2)

[1] 0.1612245
```

The **caret** package supports generalized linear models fit by `glm()` and can straightforwardly perform LOO CV:

```
R> library("caret", quietly=TRUE, warn.conflicts=FALSE)
```

Attaching package: 'ggplot2'

The following objects are masked \_by\_ '.GlobalEnv':

layer, mpg

The following object is masked from 'package:latticeExtra':

layer

```
R> # reports complement of Bayes-rule criterion:
R> train(x = attrition[, c("JobSatisfaction", "Gender",
+                           "MonthlyIncome")],
+       y = attrition$Attrition,
+       method = "glm",
+       trControl = trainControl(method="LOOCV"))
```

Generalized Linear Model

```
1470 samples
  3 predictor
  2 classes: 'No', 'Yes'
```

No pre-processing

Resampling: Leave-One-Out Cross-Validation

Summary of sample sizes: 1469, 1469, 1469, 1469, 1469, 1469, ...



```

+ },
+ caret=train(x = attrition[, c("JobSatisfaction", "Gender",
+                               "MonthlyIncome")],
+             y = attrition$Attrition,
+             method = "glm",
+             trControl = trainControl(method="LOOCV")),
+ direct={
+   n <- nrow(attrition)
+   yhat <- numeric(n)
+   for (i in 1:n){
+     m <- update(mod.attrition, data=attrition[-i, ])
+     yhat[i] <- predict(m, newdata=attrition[i, ], type="response")
+   }
+   mean(((attrition$Attrition == "Yes") - round(yhat))^2)
+ },
+ cv.exact=cv(mod.attrition, k="loo", criterion=BayesRule),
+ cv.wood=cv(mod.attrition, k="loo", criterion=BayesRule,
+            method="Woodbury"),
+ cv.hatvalues=cv(mod.attrition, k="loo", criterion=BayesRule,
+                method="hatvalues"),
+ times=10
+ )

```

Unit: milliseconds

	expr	min	lq	mean	median	uq
	rsample	5875.568755	5937.488595	6022.387836	5992.053568	6088.319477
	caret	6817.660646	6885.694078	6985.358420	6960.156002	7044.248130
	direct	5365.392799	5438.409535	5470.481137	5460.706996	5503.486744
	cv.exact	4858.216649	4867.834306	4942.216059	4894.608147	4969.956901
	cv.wood	105.610383	106.629766	109.464535	107.184250	112.978821
	cv.hatvalues	1.352672	1.483462	1.551227	1.511096	1.601747
	max neval		cld			
	6316.487921	10	a			
	7216.970962	10	b			
	5599.985287	10	c			
	5245.955330	10	d			
	119.953741	10	e			
	1.812528	10	e			

Computing time for **rsample** and **caret** for this problem is similar to direct computation of LOO CV and to `cv()` with `method="exact"`, but more than an order of magnitude slower than `cv()` with `method="Woodbury"`, and more than three orders of magnitude slower than `cv()` with `method="hatvalues"`.

## 7. Computational notes

### 7.1. Efficient computations for linear and generalized linear models

The most straightforward way to implement cross-validation in R for statistical modeling functions that are written in the canonical manner is to use `update()` to refit the model with each fold removed. This is the approach taken in the default method for `cv()`, and it is appropriate if the cases are independently sampled. Refitting the model in this manner for each fold is generally feasible when the number of folds is modest, but can be prohibitively costly for leave-one-out cross-validation when the number of cases is large.

The `"lm"` and `"glm"` methods for `cv()` take advantage of computational efficiencies by avoiding refitting the model with each fold removed. Consider, in particular, the weighted linear model  $\mathbf{y}_{n \times 1} = \mathbf{X}_{n \times p} \boldsymbol{\beta}_{p \times 1} + \boldsymbol{\varepsilon}_{n \times 1}$ , where  $\boldsymbol{\varepsilon} \sim \mathbf{N}_n(\mathbf{0}, \sigma^2 \mathbf{W}_{n \times n}^{-1})$ . Here,  $\mathbf{y}$  is the response vector,  $\mathbf{X}$  the model matrix, and  $\boldsymbol{\varepsilon}$  the error vector, each for  $n$  cases, and  $\boldsymbol{\beta}$  is the vector of  $p$  population regression coefficients. The errors are assumed to be multivariately normally distributed with 0 means and covariance matrix  $\sigma^2 \mathbf{W}^{-1}$ , where  $\mathbf{W} = \text{diag}(w_i)$  is a diagonal matrix of inverse-variance weights. For the linear model with constant error variance, the weight matrix is taken to be  $\mathbf{W} = \mathbf{I}_n$ , the order- $n$  identity matrix.

The weighted-least-squares (WLS) estimator of  $\boldsymbol{\beta}$  is (see, e.g., [Fox 2016](#), Sec. 12.2.2) <sup>14</sup>

$$\mathbf{b}_{\text{WLS}} = (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} \mathbf{y}$$

Fitted values are then  $\hat{\mathbf{y}} = \mathbf{X} \mathbf{b}_{\text{WLS}}$ .

The LOO fitted value for the  $i$ th case can be efficiently computed by  $\hat{y}_{-i} = y_i - e_i / (1 - h_i)$  where  $h_i = \mathbf{x}_i^T (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{x}_i$  (the so-called “hatvalue”). Here,  $\mathbf{x}_i^T$  is the  $i$ th row of  $\mathbf{X}$ , and  $\mathbf{x}_i$  is the  $i$ th row written as a column vector. This approach can break down when one or more hatvalues are equal to 1, in which case the formula for  $\hat{y}_{-i}$  requires division by 0. In this case, the “training” set omitting the observation with hatvalue = 1 is rank-deficient and the predictors for the left-out case are outside the linear span of the predictors in the training set.

To compute cross-validated fitted values when the folds contain more than one case, we make use of the Woodbury matrix identity ([Hager 1989](#)),

$$(\mathbf{A}_{m \times m} + \mathbf{U}_{m \times k} \mathbf{C}_{k \times k} \mathbf{V}_{k \times m})^{-1} = \mathbf{A}^{-1} - \mathbf{A}^{-1} \mathbf{U} (\mathbf{C}^{-1} + \mathbf{V} \mathbf{A}^{-1} \mathbf{U})^{-1} \mathbf{V} \mathbf{A}^{-1}$$

where  $\mathbf{A}$  is a nonsingular order- $n$  matrix. We apply this result by letting

$$\begin{aligned} \mathbf{A} &= \mathbf{X}^T \mathbf{W} \mathbf{X} \\ \mathbf{U} &= \mathbf{X}_{\mathbf{j}}^T \\ \mathbf{V} &= -\mathbf{X}_{\mathbf{j}} \\ \mathbf{C} &= \mathbf{W}_{\mathbf{j}} \end{aligned}$$

where the subscript  $\mathbf{j} = (i_{j1}, \dots, i_{jm})^T$  represents the vector of indices for the cases in the  $j$ th fold,  $j = 1, \dots, k$ . The negative sign in  $\mathbf{V} = -\mathbf{X}_{\mathbf{j}}$  reflects the *removal*, rather than addition, of the cases in  $\mathbf{j}$ .

<sup>14</sup>This is a definitional formula, which assumes that the model matrix  $\mathbf{X}$  is of full column rank, and which can be subject to numerical instability when  $\mathbf{X}$  is ill-conditioned. `lm()` uses the singular-value decomposition of the model matrix to obtain computationally more stable results.

Applying the Woodbury identity isn't quite as fast as using the hatvalues, but it is generally much faster than refitting the model. A disadvantage of the Woodbury identity, however, is that it entails explicit matrix inversion and thus may be numerically unstable. The inverse of  $\mathbf{A} = \mathbf{X}^T \mathbf{W} \mathbf{X}$  is available directly in the "lm" object, but the second term on the right-hand side of the Woodbury identity requires a matrix inversion with each fold deleted. (In contrast, the inverse of each  $\mathbf{C} = \mathbf{W}_j$  is straightforward because  $\mathbf{W}$  is diagonal.)

The Woodbury identity also requires that the model matrix be of full rank. We impose that restriction in our code by removing redundant regressors from the model matrix for all of the cases, but that doesn't preclude rank deficiency from surfacing when a fold is removed. Rank deficiency of  $\mathbf{X}$  doesn't disqualify cross-validation because all we need are fitted values under the estimated model.

`glm()` computes the maximum-likelihood estimates for a generalized linear model by iterated weighted least squares (see, e.g., [Fox and Weisberg 2019](#), Sec. 6.12). The last iteration is therefore just a WLS fit of the "working response" on the model matrix using "working weights." Both the working weights and the working response at convergence are available from the information in the object returned by `glm()`.

We then treat re-estimation of the model with a case or cases deleted as a WLS problem, using the hatvalues or the Woodbury matrix identity. The resulting fitted values for the deleted fold aren't exact—that is, except for the Gaussian family, the result isn't identical to what we would obtain by literally refitting the model—but in our (limited) experience, the approximation is very good, especially for LOO CV, which is when we would be most tempted to use it. Nevertheless, because these results are approximate, the default for the "glm" `cv()` method is to perform the exact computation, which entails refitting the model with each fold omitted.

## 7.2. Computation of the bias-corrected CV criterion and confidence intervals

Let  $CV(\mathbf{y}, \hat{\mathbf{y}})$  represent a cross-validation cost criterion, such as mean-squared error, computed for all of the  $n$  values of the response  $\mathbf{y}$  based on fitted values  $\hat{\mathbf{y}}$  from the model fit to all of the data. We require that  $CV(\mathbf{y}, \hat{\mathbf{y}})$  is the mean of casewise components, that is,  $CV(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{n} \sum_{i=1}^n cv(y_i, \hat{y}_i)$ .<sup>15</sup> For example,  $MSE(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$ .<sup>16</sup>

We divide the  $n$  cases into  $k$  folds of approximately  $n_j \approx n/k$  cases each, where  $n = \sum n_j$ . As above, let  $\mathbf{j}$  denote the indices of the cases in the  $j$ th fold.

Now define  $CV_j = CV(\mathbf{y}, \hat{\mathbf{y}}^{(j)})$ . The superscript  $(j)$  on  $\hat{\mathbf{y}}^{(j)}$  represents fitted values computed for all of the cases from the model with fold  $j$  omitted. Let  $\hat{\mathbf{y}}^{(-i)}$  represent the vector of fitted values for all  $n$  cases where the fitted value for the  $i$ th case is computed from the model fit with the fold including the  $i$ th case omitted (i.e., fold  $j$  for which  $i \in \mathbf{j}$ ).

Then the cross-validation criterion is just  $CV = CV(\mathbf{y}, \hat{\mathbf{y}}^{(-i)})$ . Following [Davison and Hinkley](#)

<sup>15</sup>[Arlot and Celisse \(2010\)](#) term the casewise loss,  $cv(y_i, \hat{y}_i)$ , the "contrast function."

<sup>16</sup>Some commonly employed CV criteria—such as the root-mean-squared error (RMSE), median absolute error, and, for binary-regression models, the complement of the area under the receiver operating characteristic (ROC) curve—are not means of casewise components. That the RMSE and median absolute error aren't means of casewise components is obvious; for the complement of the area under the ROC curve, see the vignette on extending the `cv` package.



(1997, pp. 293–295), the bias-adjusted cross-validation criterion is

$$CV_{\text{adj}} = CV + CV(\mathbf{y}, \hat{\mathbf{y}}) - \frac{1}{n} \sum_{j=1}^k n_j CV_j$$

We compute the standard error of CV as

$$SE(CV) = \frac{1}{\sqrt{n}} \sqrt{\frac{\sum_{i=1}^n [cv(y_i, \hat{y}_i^{(-i)}) - CV]^2}{n-1}}$$

that is, as the standard deviation of the casewise components of CV divided by the square-root of the number of cases.

We then use  $SE(CV)$  to construct a  $100 \times (1 - \alpha)\%$  confidence interval around the *adjusted* CV estimate of error:

$$[CV_{\text{adj}} - z_{1-\alpha/2} SE(CV), CV_{\text{adj}} + z_{1-\alpha/2} SE(CV)]$$

where  $z_{1-\alpha/2}$  is the  $1 - \alpha/2$  quantile of the standard-normal distribution (e.g,  $z \approx 1.96$  for a 95% confidence interval, for which  $1 - \alpha/2 = .975$ ).

Bates *et al.* (2023) show that the coverage of this confidence interval is poor for small samples, and they suggest a much more computationally intensive procedure, called *nested cross-validation*, to compute better estimates of error and confidence intervals with better coverage for small samples. We may implement Bates et al.’s approach in a later release of the **cv** package. At present we use the confidence interval above for sufficiently large  $n$ , which, based on Bates et al.’s results, we take by default to be  $n \geq 400$ .

## References

- Andersen R, Heath A, Sinnott R (2002). “Political knowledge and electoral choice.” *British Elections and Parties Review*, **12**, 11–27.
- Arlot S, Celisse A (2010). “A survey of cross-validation procedures for model selection.” *Statistics Surveys*, **4**, 40 – 79. URL <https://doi.org/10.1214/09-SS054>.
- Barnard GA (????). “Discussion of Professor Stone’s paper.” *Journal of the Royal Statistical Society, Series B*, **36**, 133–135.
- Bates D, Mächler M, Bolker B, Walker S (2015). “Fitting Linear Mixed-Effects Models Using lme4.” *Journal of Statistical Software*, **67**(1), 1–48.
- Bates S, Hastie T, Tibshirani R (2023). “Cross-validation: What does it estimate and how well does it do it?” *Journal of the American Statistical Association*, **in press**. URL <https://doi.org/10.1080/01621459.2023.2197686>.
- Box GEP, Cox DR (1964). “An analysis of transformations.” *Journal of the Royal Statistical Society, Series B*, **26**, 211–252.

- Brooks ME, Kristensen K, van Benthem KJ, Magnusson A, Berg CW, Nielsen A, Skaug HJ, Maechler M, Bolker BM (2017). “glmmTMB Balances Speed and Flexibility Among Packages for Zero-inflated Generalized Linear Mixed Modeling.” *The R Journal*, **9**(2), 378–400. doi:10.32614/RJ-2017-066.
- Canty A, Ripley BD (2022). *boot: Bootstrap R (S-Plus) Functions*. R package version 1.3-28.1.
- Coyle J, Hejazi N, Malenica I, Phillips R (2022). *origami: Generalized Framework for Cross-Validation*. R package version 1.0.7, URL <https://CRAN.R-project.org/package=origami>.
- Davison AC, Hinkley DV (1997). *Bootstrap Methods and Their Applications*. Cambridge University Press, Cambridge.
- Fox J (2016). *Applied Regression Analysis and Generalized Linear Models*. Second edition edition. Sage, Thousand Oaks CA.
- Fox J, Weisberg S (2019). *An R Companion to Applied Regression*. Third edition edition. Sage, Thousand Oaks CA.
- Frick H, Chow F, Kuhn M, Mahoney M, Silge J, Wickham H (2024). *rsample: General Resampling Infrastructure*. R package version 1.2.1, URL <https://CRAN.R-project.org/package=rsample>.
- Hager WW (1989). “Updating the inverse of a matrix.” *SIAM Review*, **31**(2), 221–239.
- Harrell Jr F (2015). *Regression Modeling Strategies*. Second edition edition. Springer, New York.
- Hastie T, Tibshirani R, Friedman J (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Second edition edition. Springer, New York. URL [https://hastie.su.domains/ElemStatLearn/printings/ESLII\\_print12\\_toc.pdf](https://hastie.su.domains/ElemStatLearn/printings/ESLII_print12_toc.pdf).
- James G, Witten D, Hastie T, Tibshirani R (2021). *An Introduction to Statistical Learning with Applications in R*. Second edition edition. Springer, New York.
- Kapsner LA (2024). *ml experiments: Machine Learning Experiments*. R package version 0.0.3, URL <https://CRAN.R-project.org/package=ml experiments>.
- Kuhn M (2008). “Building Predictive Models in R Using the caret Package.” *Journal of Statistical Software*, **28**(5), 1–26. doi:10.18637/jss.v028.i05. URL <https://www.jstatsoft.org/index.php/jss/article/view/v028i05>.
- Kuhn M (2024). *modeldata: Data Sets Useful for Modeling Examples*. R package version 1.3.0, URL <https://CRAN.R-project.org/package=modeldata>.
- Lüdecke D, Waggoner P, Makowski D (2019). “insight: A Unified Interface to Access Information from Model Objects in R.” *Journal of Open Source Software*, **4**(38), 1412.
- Mersmann O (2023). *microbenchmark: Accurate Timing Functions*. R package version 1.4.10, URL <https://CRAN.R-project.org/package=microbenchmark>.

- Olsen LR, Zachariae HB (2024). *cvms: Cross-Validation for Model Selection*. R package version 1.6.1, URL <https://CRAN.R-project.org/package=cvms>.
- Pfizinger J (2024). *tidyfit: Regularized Linear Modeling with Tidy Data*. R package version 0.7.1, URL <https://CRAN.R-project.org/package=tidyfit>.
- Pinheiro JC, Bates DM (2000). *Mixed-Effects Models in S and S-PLUS*. Springer, New York.
- Raudenbush SW, Bryk AS (2002). *Hierarchical Linear Models: Applications and data analysis methods*. Second edition edition. Sage, Thousand Oaks CA.
- Sarkar D (2008). *Lattice: Multivariate Data Visualization with R*. Springer, New York. ISBN 978-0-387-75968-5. URL <http://lmdvr.r-forge.r-project.org>.
- Sarkar D, Andrews F (2022). *latticeExtra: Extra Graphical Utilities Based on lattice*. R package version 0.6-30, URL <https://CRAN.R-project.org/package=latticeExtra>.
- Vehtari A (2023). “Cross-validation FAQ.” URL <https://users.aalto.fi/~ave/CV-FAQ.html>.
- Venables WN, Ripley BD (2002). *Modern Applied Statistics with S*. Fourth edition edition. Springer, New York.
- Weisberg S (2014). *Applied Linear Regression*. Second edition edition. Wiley, Hoboken NJ.

### Affiliation:

John Fox  
 McMaster University  
 Hamilton, Ontario, Canada  
 E-mail: [jfox@mcmaster.ca](mailto:jfox@mcmaster.ca)  
 URL: <https://www.john-fox.ca/>

Georges Monette  
 York University  
 Toronto, Ontario, Canada